

DBMSI VIVA

Viva Questions & Answers – DMSL Practicals (1–5)

Practical 1 – Basic SQL Table Creation and Constraints

Q1. What is a primary key?

A. A column that uniquely identifies each record in a table.

Q2. What is a foreign key?

A. A key used to link two tables; it references the primary key of another table.

Q3. Can a table have multiple primary keys?

A. No, only one primary key per table, but it can be composite.

Q4. What is a default constraint?

A. It assigns a predefined value to a column when no value is given.

Q5. What is the purpose of a check constraint?

A. It ensures that all values in a column meet a specific condition.

Practical 2 – PL/SQL Block (Control Structure & Exception Handling)

Q1. What are the sections in a PL/SQL block?

A. DECLARE, BEGIN, EXCEPTION, and END.

Q2. What is an exception in PL/SQL?

A. An error or abnormal condition that disrupts normal program flow.

Q3. What are the two types of exceptions?

A. Predefined and user-defined exceptions.

Q4. What is the use of the DECLARE section?

A. It is used to declare variables and constants.

Q5. How does PL/SQL handle exceptions?

A. Using the EXCEPTION section with handlers like `WHEN NO_DATA_FOUND THEN`.

Practical 3 – Customer & Account (SQL Joins & Aggregation)

Q1. What is a join in SQL?

A. It combines data from two or more tables based on a related column.

Q2. Difference between INNER JOIN and LEFT JOIN?

A. INNER JOIN shows only matching rows; LEFT JOIN shows all rows from the left table.

Q3. What is a natural join?

A. A join that automatically matches columns with the same name in both tables.

Q4. What does GROUP BY do?

A. Groups rows sharing the same values so aggregate functions can be applied.

Q5. What is the HAVING clause used for?

A. To filter groups created by GROUP BY after aggregation.

Practical 4 – Library Fine Calculation (PL/SQL Block)

Q1. What is the purpose of SYSDATE?

A. It returns the current system date and time.

Q2. What is a cursor in PL/SQL?

A. A pointer that allows row-by-row processing of query results.

Q3. How do you calculate fine in PL/SQL?

A. By comparing the due date and return date, applying a rate per late day.

Q4. What is an IF-ELSE statement used for?

A. For conditional execution of statements based on logical tests.

Q5. Why is status updated from 'I' to 'R'?

A. To indicate that the book is returned instead of issued.

Practical 5 – Student–Instructor–Course Database

Q1. What does the HAVING clause do?

A. It filters aggregated results in a GROUP BY query.

Q2. How can we increase salary by 10% in SQL?

A. Using `UPDATE table SET salary = salary * 1.10`.

Q3. What is the use of NOT IN operator?

A. To exclude specific values from a result set.

Q4. How do you find substrings in names?

A. By using `LIKE '%am%'` in the WHERE clause.

Q5. What is the purpose of a JOIN in this practical?

A. To combine student and course data to show which student takes which subject.

Viva Questions & Answers – DMSL Practicals (6–10)

Practical 6 – Cursors in PL/SQL

Q1. What is a cursor in PL/SQL?

A. A cursor is a pointer to the result set of a query that allows row-by-row processing.

Q2. What are the types of cursors?

A. Implicit (handled automatically) and explicit (defined by the programmer).

Q3. What are cursor attributes?

A. `%FOUND`, `%NOTFOUND`, `%ROWCOUNT`, and `%ISOPEN` give cursor status information.

Q4. What is a parameterized cursor?

A. A cursor that accepts parameters to filter data dynamically.

Q5. Why use cursors in PL/SQL?

A. To handle multiple rows returned by a query one at a time.

Practical 7 – Stored Procedure (Grade Calculation)

Q1. What is a stored procedure?

A. A precompiled block of SQL statements stored in the database.

Q2. What are IN, OUT, and INOUT parameters?

A. IN passes data to the procedure, OUT returns data, and INOUT does both.

Q3. How is a procedure different from a function?

A. A procedure doesn't return a value directly; a function always returns one.

Q4. What are the advantages of stored procedures?

A. Better performance, reusability, and centralized logic.

Q5. How do you execute a procedure?

A. Using the `CALL procedure_name();` statement in MySQL.

Practical 8 – PL/SQL Loop for Circle Area

Q1. What are the types of loops in PL/SQL?

A. Simple loop, FOR loop, and WHILE loop.

Q2. What is the difference between FOR and WHILE loops?

A. FOR has a fixed range; WHILE runs until a condition becomes false.

Q3. How can you exit a loop early?

A. Using the `EXIT` or `EXIT WHEN` statement.

Q4. What is the purpose of using loops here?

A. To repeatedly calculate and insert multiple records like circle areas.

Q5. What is the syntax of a FOR loop?

A. `FOR i IN start..end LOOP ... END LOOP;`

Practical 9 – Triggers

Q1. What is a trigger?

A. A stored block that executes automatically when a specific event occurs on a table.

Q2. What are types of triggers?

A. BEFORE, AFTER, and INSTEAD OF triggers.

Q3. When does a BEFORE trigger execute?

A. Before an insert, update, or delete operation occurs.

Q4. What is the practical use of triggers?

A. For enforcing business rules or automatic logging.

Q5. Can a trigger call another trigger?

A. Yes, if an operation inside it causes another triggering event.

Practical 10 – MongoDB OrderInfo Collection

Q1. What is a collection in MongoDB?

A. A collection is similar to a table in SQL that stores multiple JSON-like documents.

Q2. What is a document?

A. A single record in MongoDB represented in BSON (Binary JSON) format.

Q3. What is the use of `$match` in aggregation?

A. It filters documents based on given conditions.

Q4. How do you create an index in MongoDB?

A. Using `db.collection.createIndex({ field: 1 })`.

Q5. What is the difference between `insertOne()` and `insertMany()`?

A. `insertOne()` adds a single document, while `insertMany()` adds multiple at once.

Viva Questions & Answers – DMSL Practicals (11-15)

Practical 11 – MongoDB Movies Collection

Q1. What does `insertOne()` do?

A. It inserts a single document into a MongoDB collection.

Q2. How do you find documents with a condition?

A. Using `find()` with a filter, e.g. `db.movies.find({budget:{$gt:100000}})`.

Q3. What does `$gt` mean in MongoDB?

A. It means "greater than" and filters values above a given number.

Q4. How do you update multiple documents?

A. Using `updateMany()` instead of `updateOne()`.

Q5. What is dot notation used for?

A. To access nested fields inside embedded documents, like `"producer.name"`.

Practical 12 – MongoDB Teachers–Department–Students

Q1. What is a nested document in MongoDB?

A. A document that contains another document as a field value.

Q2. How do you use `$or` in MongoDB?

A. It filters documents matching at least one of multiple conditions.

Q3. What is the use of `updateOne()`?

A. It updates the first document that matches a filter.

Q4. How do you delete multiple documents?

A. Using `deleteMany()` with a condition.

Q5. Why create indexes in MongoDB?

A. To improve query performance and speed of searching.

Practical 13 – MongoDB Aggregation and Indexing

Q1. What is aggregation in MongoDB?

A. A process of performing calculations on grouped data like SQL GROUP BY.

Q2. What does `$group` do?

A. Groups documents by a field and applies accumulators like `$sum` , `$avg` .

Q3. What does `$avg` calculate?

A. It returns the average value of a numeric field.

Q4. How do you drop an index?

A. Using `db.collection.dropIndex("index_name")` .

Q5. Why use aggregation instead of simple queries?

A. To perform complex computations like totals and averages in one query.

Practical 14 – MongoDB MapReduce

Q1. What is MapReduce?

A. It's a data processing model that maps data to key-value pairs and reduces results.

Q2. What is the role of the map function?

A. It emits key-value pairs for processing.

Q3. What is the role of the reduce function?

A. It combines and summarizes values with the same key.

Q4. What does `emit()` do in MapReduce?

A. Emits a key and its value during the map phase.

Q5. Difference between aggregation and MapReduce?

A. Aggregation is simpler and faster; MapReduce is more flexible for large data sets.

Practical 15 – MongoDB Student Marks Aggregation

Q1. What does `$max` operator do?

A. It returns the highest value from grouped documents.

Q2. What does `$min` operator do?

A. It returns the lowest value from grouped documents.

Q3. What does `$sum` calculate?

A. The total of numeric fields in a group.

Q4. What does `$avg` represent?

A. The average value of numeric data.

Q5. What is the purpose of `$group` in aggregation?

A. To group documents based on a specific field for summarization.

Viva Questions & Answers – DMSL Practicals (16-20)

Practical 16 – MongoDB MapReduce (Salary Calculation)

Q1. What is the purpose of the map function?

A. It processes each document and emits key-value pairs for reduction.

Q2. What is the reduce function used for?

A. It aggregates and combines all values that share the same key.

Q3. What is MapReduce mainly used for?

A. For complex calculations like total salary or count across large data sets.

Q4. What is the use of the `out` option in MapReduce?

A. It defines where the output of MapReduce will be stored.

Q5. Difference between MapReduce and aggregation?

A. Aggregation is faster; MapReduce is more flexible and customizable.

Practical 17 – MongoDB OrderInfo Updates

Q1. How do you remove a field from a document?

A. Using `$unset`, e.g., `db.collection.update({}, {$unset:{field:1}})`.

Q2. What does `$or` do?

A. Combines multiple conditions, returning documents that match any.

Q3. How do you increment a value?

A. Using `$inc`, e.g., `db.collection.update({}, {$inc:{amount:100}})`.

Q4. Difference between `$inc` and `$mul` ?

A. `$inc` adds a value, `$mul` multiplies a value.

Q5. What is `$set` used for?

A. To update or add a specific field in a document.

Practical 18 – SQL Employee–Company Database

Q1. What is a `CASE` statement used for?

A. To apply conditional logic within a SQL statement.

Q2. What is the difference between `UPDATE` and `ALTER` ?

A. `UPDATE` changes data; `ALTER` modifies table structure.

Q3. What is a foreign key?

A. It enforces a relationship between two tables.

Q4. What is the use of `DELETE` ?

A. To remove records from a table while keeping the structure.

Q5. What is a composite key?

A. A key formed by combining two or more columns to uniquely identify a record.

Practical 19 – MongoDB Employee Collection

Q1. What is the difference between `insertOne()` and `insertMany()` ?

A. `insertOne()` adds a single document; `insertMany()` adds multiple at once.

Q2. How do you drop a collection?

A. Using `db.collection.drop()` .

Q3. How do you rename a collection?

A. Using `db.collection.renameCollection("newName")` .

Q4. What is `$regex` used for?

A. For pattern matching in string searches.

Q5. What does `sort()` do in MongoDB?

- A. It orders the output by specified fields, ascending or descending.
-

Practical 20 – SQL Duty Allocation

Q1. What is a view in SQL?

- A. A virtual table created using the result of a SELECT query.

Q2. What is an aggregate function?

- A. A function that performs a calculation on multiple values, e.g., `SUM()`, `AVG()`.

Q3. What does `GROUP BY` do?

- A. Groups rows with the same values for aggregation.

Q4. What is a foreign key used for?

- A. To maintain referential integrity between two tables.

Q5. What is the purpose of `COUNT()`?

- A. To return the total number of rows that meet a condition.

Viva Questions & Answers – DMSL Practicals (21–23)

Practical 21 – Employee–Company–Manager (SQL Joins & Updates)

Q1. What is a self-join?

- A. It's a join where a table is joined with itself to compare related data in the same table.

Q2. What does `NOT IN` do?

- A. It excludes rows that match a list or subquery of specified values.

Q3. What is the purpose of aliasing (`AS`)?

- A. To assign a temporary name to a table or column for readability.

Q4. How do you update salaries in SQL?

- A. Using the `UPDATE` statement with arithmetic operations like `salary = salary * 1.10`.

Q5. What is the difference between `DELETE` and `DROP` ?

A. `DELETE` removes records, while `DROP` removes the entire table structure.

Practical 22 – Advanced Employee–Company Queries

Q1. What is a correlated subquery?

A. A subquery that depends on values from the outer query to execute.

Q2. What does `> ALL` mean in SQL?

A. It compares a value to all values in a subquery and returns true if greater than every one.

Q3. What is the use of `NOT EXISTS` ?

A. It returns rows when the subquery finds no matching rows.

Q4. What is the role of `HAVING` in aggregation?

A. It filters aggregated data after a `GROUP BY` operation.

Q5. How can you find the company with the maximum employees?

A. Use `COUNT()` with `GROUP BY` and `ORDER BY DESC LIMIT 1`.

Practical 23 – Employee Table (Constraints, Index & View)

Q1. What is an index in SQL?

A. A data structure that improves query performance by speeding up searches.

Q2. What is `AUTO_INCREMENT` used for?

A. To automatically generate sequential numeric values for a column.

Q3. What is the difference between a table and a view?

A. A table stores physical data, while a view is a virtual table derived from a query.

Q4. What does the `DEFAULT` constraint do?

A. It automatically assigns a predefined value if no value is given during insert.

Q5. How do you add a new column after table creation?

A. Using `ALTER TABLE table_name ADD column_name datatype;` .

Create following tables with primary and foreign key and solve the queries given below

Person (driver_id, name, address)

Car (license, model, year)

Accident (report_no, date_acc, location)

Owns (driver_id, license)

Participated (driver_id, model, report_no, damage_amount)



3. Views

Q9. What is a view?

A: A view is a virtual table based on the result of an SQL query.

Q10. Can we update data through a view?

A: Yes, if the view is based on a single table and doesn't use aggregation or joins that prevent updates.

Q11. How did you create `Driver_Car_View` ?

A: By joining `Person`, `Owns`, and `Car` tables to show each driver's name and car model.

Q12. What is the difference between `CREATE VIEW` and `CREATE OR REPLACE VIEW` ?

A: `CREATE` is used initially, while `CREATE OR REPLACE` modifies an existing view without dropping it first.

12
34

4. Sequences

Q13. What is a sequence in SQL?

A: A sequence automatically generates numeric values, often used for primary keys.

Q14. What is the starting value and increment in `seq_report_no` ?

A: It starts at **6001** and increments by **5**.

Q15. How do you use a sequence in an `INSERT` statement?

A: `INSERT INTO Person VALUES (seq_driver_id.NEXTVAL, 'Name', 'City');`



5. Indexes

Q16. What is an index?

A: An index improves the speed of data retrieval operations.

Q17. What is the difference between a normal index and a unique index?

A: A **unique index** ensures all values in the indexed column are distinct.

Q18. What is a composite index?

A: An index created on multiple columns, e.g., `CREATE INDEX idx_car_model_year ON Car(model, year);`

Q19. Why did we create an index on `damage_amount` ?

A: To speed up queries that frequently filter or sort by `damage_amount` .



6. Synonyms

Q20. What is a synonym in SQL?

A: A synonym is an alias for a database object like a table or view.

Q21. Why do we create a synonym for the `Accident` table?

A: To simplify queries — instead of typing `Accident` , we can use `acc_info` .

Q22. Can synonyms be public or private?

A: Yes — **private synonyms** are accessible to one user, **public synonyms** to all users.
